


星威华商论坛 Post approval flow

How a forum post gets checked before it goes live · built on the client's 敏感词库 word list, received 4 Aug 2026
219 active keywords + 8 pattern rules · nothing publishes without passing review

- Automated — no human time
- Human moderator
- Published



ONE DECISION FOR CYS
`MID_RISK_ACTION`

You asked for high **and mid** risk to be blocked outright. High risk — agreed, straight rejection. For **mid risk we recommend the review queue instead**, and the reason is in your own file: every mid-risk row is annotated 需结合上下文判断 — *needs context to judge*. 换汇 sits on that list and it is your own core service, so 我们提供合规换汇服务 would be auto-rejected and the business would never learn why. Same for 发票, 佣金, 代付. The queue keeps the identical guarantee — still checked before it goes live — without losing good posts. Ships as `MID_RISK_ACTION`, switchable either way, so nothing waits on your answer.

What the author sees	审核状态
审核中	Under review — visible in 我的帖子, still editable
已发布	Published, with its public link
未通过	Not approved — category + rewrite hint, editable, resubmits to the queue

Every state change sends an email. Silence after submitting is the biggest risk in a pre-publication design — it reads as a broken site.

Why the word list runs first	STAGE 1 → 2
The scan is instant, free, and legally legible — "matched 逃汇, 《外汇管理条例》" is a defensible answer to a regulator; a model's opinion is not. It also means high-risk posts are rejected before any AI call is paid for.	
The AI stage exists for the word list's blind spot: negation — 反洗钱 contains 洗钱 — and posts that are sensitive without using any listed word.	

Cost and speed	PER POST
Word-list scan	~1 ms · free
AI context check	~2 s · ~\$0.0015
Running cost at 50 posts/day	~\$2 / month
Human review	flagged posts only

The word list keeps the AI bill near zero by rejecting the worst content before it is ever sent.